

Program analysis application: memory safety

Gang Tan

CSE 597: Special topics on program analysis

Spring 2021

Penn State University

Why study this?

- Memory safety is an important program analysis application domain
 - Numerous dynamic/static analysis studied for identifying memory safety violations in C/C++ applications
- Some knowledge here needed for some of our homework (e.g., fuzzing)

Some Terminology

- Software **error**
 - A programming mistake that makes the software not meet its expectation
- Software **vulnerability**
 - A software error that can lead to possible attacks
- **Attack**
 - The process of exploiting a vulnerability
 - An attack can exploit a vulnerability to achieve additional functionalities for attackers
 - E.g., privilege escalation, arbitrary code execution

Language of Choice for Systems Programming: C/C++

- Systems software
 - OS; hypervisor; web servers; firmware; network controllers; device drivers; compilers; ...
- Benefits of C/C++: programming model close to the machine model; flexible; efficient
- **BUT error-prone**
 - C/C++ not memory safe; huge security risk
 - Debugging memory errors is a headache
 - Perhaps on par with debugging multithreaded programs

Buffer Overflows

- Refer to reading/writing a buffer out of its bounds
 - Programmers' job in C/C++ to not do this
 - In contrast, many modern languages (Java, Python, ...) prevent buffer overflows by performing automatic bounds checking
- The first Internet worm, and many subsequent ones (CodeRed, Blaster, ...), exploited buffer overflows
- Buffer overflows still cause many security alerts nowadays
 - E.g., check out CERT, cve.mitre.org, or bugtraq

C STRINGS: USAGE AND PITFALLS

C-Style Strings

- C-style strings consist of a contiguous sequence of characters, terminated by and including the first null character.
 - String length is the number of bytes preceding the null character.
 - The number of bytes required to store a string is the number of characters plus one (times the size of each character).

h	e	l	l	o	\0
---	---	---	---	---	----

Using Strings in C

- C provides many string functions in its libraries (libc)
- For example, we use the strcpy function to copy one string to another:

```
#include <string.h>
char string1[] = "Hello, world!";
char string2[20];
strcpy(string2, string1);
```


Using Strings in C

- Another lets us compare strings

```
char string3[] = "this is";  
char string4[] = "a test";  
if(strcmp(string3, string4) == 0)  
    printf("strings are equal\n");  
else printf("strings are different\n")
```

- This code fragment will print "strings are different". Notice that strcmp does **not** return a boolean result.

Other Common String Functions

- strlen: getting the length of a string
- strcpy/strncpy: string copying
- strcat/strncat: string concatenation
- gets, fgets: receive input to a string
- ...

Common String Manipulation Errors

- Programming with C-style strings, in C or C++, is error prone
- Common errors include
 - Buffer overflows
 - null-termination errors
 - off-by-one errors
 - ...

gets: Unbounded String Copies

- Occur when data is copied from an unbounded source to a fixed-length character array

```
void main(void) {  
    char Password[8];  
    puts("Enter a 8-character password:");  
    gets(Password);  
    printf("Password=%s\n", Password);  
}
```

strcpy and strcat

- The standard string library functions do not know the size of the destination buffer

```
int main(int argc, char *argv[]) {  
    char name[2048];  
    strcpy(name, argv[1]);  
    strcat(name, " = ");  
    strcat(name, argv[2]);  
    ...  
}
```

Better String Library Functions

- Functions that restrict the number of bytes are often recommended
- Never use `gets(buf)`
 - Use `fgets(buf, size, stdin)` instead

From gets to fgets

- `char *fgets(char *BUF, int N, FILE *FP);`
 - *“Reads at most N-1 characters from FP until a newline is found. The characters including to the newline are stored in BUF. The buffer is terminated with a 0.”*

```
void main(void) {  
    char Password[8];  
    puts("Enter a 8-character password:");  
    fgets(Password, 8, stdin);  
    ...  
}
```

Better String Library Functions

- Instead of `strcpy()`, use `strncpy()`
- Instead of `strcat()`, use `strncat()`
- Instead of `sprintf()`, use `snprintf()`

But Still Need Care

- `char *strncpy(char *s1, const char *s2, size_t n);`
 - *“Copy not more than n characters (including the null character) from the array pointed to by $s2$ to the array pointed to by $s1$; If the string pointed to by $s2$ is shorter than n characters, null characters are appended to the destination array until a total of n characters have been written.”*
 - What happens if the size of $s2$ is n or greater
 - It gets truncated
 - **And $s1$ may not be null-terminated!**

Null-Termination Errors

```
int main(int argc, char* argv[]) {  
    char a[16], b[16];  
    strncpy(a, "0123456789abcdef", sizeof(a));  
    printf("%s\n", a);  
    strcpy(b, a);  
}
```

a[] not properly terminated. Possible segmentation fault if `printf("%s\n", a);`

How to fix it?

strcpy to strncpy

- Don't replace

`strcpy(dest, src)`

by

`strncpy(dest, src, sizeof(dest))`

but by

`strncpy(dest, src, sizeof(dest)-1)`

`dst[sizeof(dest)-1] = '\0';`

if dest should be null-terminated!

- You never have this headache in memory-safe languages
- Further, strncpy has big performance penalty vs. strcpy
 - It NUL-fills the remainder of the destination

Signed vs Unsigned Numbers

```
char buf[N];  
int i, len;
```

We forget to check for negative lengths

```
read(fd, &len, sizeof(len));  
if (len > N)  
    {error ("invalid length"); return; }  
read(fd, buf, len);
```

len cast to unsigned and negative length overflows

*slide by Eric Poll

Checking for Negative Lengths

```
char buf[N];  
int i, len;  
  
read(fd, &len, sizeof(len));  
if (len > N || len < 0)  
    {error ("invalid length"); return; }  
read(fd, buf, len);
```

It still has a problem
if the buf is going to be treated as a C string.

*slide by Eric Poll

A Good Version

```
char buf[N];  
int i, len;  
  
read(fd, &len, sizeof(len));  
if (len > N-1 || len < 0)  
    {error ("invalid length"); return; }  
read(fd, buf, len);  
buf[len] = '\0'; // null terminate buf
```

*slide by Eric Poll

Integer Overflow Vulnerabilities

* slides adapted from those by Seacord

Integer Overflows

- An **integer overflow** occurs when an integer is increased beyond its maximum value or decreased beyond its minimum value
- Standard integer types (signed)
 - signed char, short int, int, long int, long long int
- Signed overflow vs unsigned overflow
 - An unsigned overflow occurs when the underlying representation can no longer represent an integer value.
 - A signed overflow occurs when a value is carried over to the sign bit

Overflow Examples

```
unsigned int ui;
```

```
signed int si;
```

```
ui = UINT_MAX; // 4,294,967,295;
```

```
ui++;
```

```
printf("ui = %u\n", ui);
```

ui = 0

```
si = INT_MAX; // 2,147,483,647
```

```
si++;
```

```
printf("si = %d\n", si);
```

si = -2,147,483,648

Overflow Examples, cont'd

```
ui = 0;
```

```
ui--;
```

```
printf("ui = %u\n", ui);
```

ui = 4,294,967,295

```
si = INT_MIN; // -2,147,483,648;
```

```
si--;
```

```
printf("si = %d\n", si);
```

si = 2,147,483,647

Integer Overflow Example

```
int main(int argc, char *const *argv) {  
    unsigned short int total;  
    total = strlen(argv[1]) + strlen(argv[2]) + 1;  
    char *buff = (char *) malloc(total);  
    strcpy(buff, argv[1]);  
    strcat(buff, argv[2]);  
}
```

What if the total variable is overflowed because of the addition operation?

Vulnerability: JPEG Example

- Based on a real-world vulnerability in the handling of the comment field in JPEG files

```
void getComment(unsigned int len, char *src) {  
    unsigned int size;  
    size = len - 2;  
    char *comment = (char *)malloc(size + 1);  
    memcpy(comment, src, size);  
    return;  
}
```

size is interpreted as a large positive value of 0xffffffff

size+1 is 0

What if I do “getComment(1, "Comment ");”?

Possible to cause an overflow by creating an image with a comment length field of 1

Vulnerability: Truncation Errors

```
int func(char *name, long cbBuf) {  
    unsigned short bufSize = cbBuf;  
    char *buf = (char *)malloc(bufSize);  
    if (buf) {  
        memcpy(buf, name, cbBuf);  
        ...  
        free(buf);  
        return 0;  
    }  
    return 1;  
}
```

What if we call the function with cbBuf greater than $2^{16} - 1$?

Heap Overflow

*adapted from slides by Trent Jaeger

Heap Overflows

- Stack overflow: overflow a memory region on the stack
- Heap overflow: overflow a memory region dynamically allocated on the heap

```
int authenticated = 0;  
char *packet = (char *)malloc(1000);
```

```
while (!authenticated) {  
    PacketRead(packet);  
    if (Authenticate(packet))  
        authenticated = 1;  
}  
if (authenticated)  
    ProcessPacket(packet);
```

What happens if PacketRead overflows the packet buffer and overwrite important data in memory?

Overflowing Heap Critical User Data

```
/* record type to allocate on heap */
typedef struct chunk {
    char inp[64];          /* vulnerable input buffer */
    void (*process)(char *); /* pointer to function */
} chunk t;

void showlen(char *buf) {
    int len; len = strlen(buf);
    printf("buffer5 read %d chars\n", len);
}

int main(int argc, char *argv[]) {
    chunk t *next;
    setbuf(stdin, NULL);
    next = malloc(sizeof(chunk t));
    next->process = showlen;
    printf("Enter value: ");
    gets(next->inp);
    next->process(next->inp);
    printf("buffer5 done\n");
}
```

example by Stallings

- Overflow the buffer on the heap so that the function pointer is changed to an arbitrary address

Overflow Heap Meta-Data


- Heap allocators (AKA memory managers)
 - What regions have been allocated and their sizes
 - What regions are available for allocation
- Heap allocators maintain metadata such as chunk size, previous, and next pointers
 - Metadata adjusted during heap-management functions
 - malloc() and free()
 - Heap metadata often adjacent to heap user data

Example Heap Allocator


- Maintain a doubly-linked list of allocated and free chunks
- `malloc()` and `free()` modify this list

An Example of Removing a Chunk


Chunks1, 2, and 3 are joined by a doubly-linked list



Chunk2 may be unlinked by rewriting 2 pointers



Chunk2 is now unlinked



- `free()` removes a chunk from allocated list
 - `chunk2->bk->fd = chunk2->fd`
 - `chunk2->fd->bk = chunk2->bk`

Attacking the Example Heap Allocator

- By overflowing chunk2, attacker controls **bk** and **fd** of chunk2
- Suppose the attacker wants to write **value** to memory address **addr**
 - Attacker sets chunk2->**fd** to be **value**
 - Attacker sets chunk2->**bk** to be **addr-offset**, where offset is the offset of the fd field in the structure

Attacking the Example Heap Allocator

- free() changed in the following way
 - `chunk2->bk->fd = chunk2->fd` becomes
`(addr-offset)->fd = value, the same as (*addr)=value`
 - `chunk2->fd->bk = chunk2->bk` becomes
`value->bk = addr-offset`
- The first memory write achieves the attacker's goal
 - Arbitrary memory writes

Use After Free and Double Free

*adapted from slides by Trent Jaeger

Use After Free

- **Error**: Program frees memory on the heap, but then references that memory as if it were still valid
 - Adversary can control data written using the freed pointer
- AKA use of dangling pointers

Use After Free

```
int main(int argc, char **argv) {  
    char *buf1, *buf2, *buf3;  
  
    buf1 = (char *) malloc(BUFSIZE1);  
  
    free(buf1);  
  
    buf2 = (char *) malloc(BUFSIZE2);  
    buf3 = (char *) malloc(BUFSIZE2);  
    strncpy(buf1, argv[1], BUFSIZE1-1);  
    ...  
}
```

What happens here?

Use After Free

- When the first buffer is freed, that memory is available for reuse right away
- Then, the following buffers are possibly allocated within that memory region

```
buf2 = (char *) malloc(BUFSIZE2);
```

```
buf3 = (char *) malloc(BUFSIZE2);
```

- Finally, the write using the freed pointer may overwrite buf2 and buf3 (and their metadata)

```
strncpy(buf1, argv[1], BUFSIZE1-1);
```

Use After Free

- Most effective attacks exploit data of another type

```
struct A {  
    void (*fnptr)(char *arg);  
    char *buf;  
};
```

```
struct B {  
    int B1;  
    int B2;  
    char info[32];  
};
```

Use After Free

- Free A, and allocate B does what?

```
x = (struct A *)malloc(sizeof(struct A));  
free(x);  
y = (struct B *)malloc(sizeof(struct B));
```

Use After Free

- How can you exploit it?

```
x = (struct A *)malloc(sizeof(struct A));  
free(x);  
y = (struct B *)malloc(sizeof(struct B));
```

```
y->B1 = 0xDEADBEEF;  
x->fnptr(x->buf);
```

- Assume that
 - The attacker controls what to write to y->B1
 - There is a later use-after-free that performs a call using “x->fnptr”
- Become a popular vulnerability to exploit – over 60% of CVEs in 2018

Exercise: Find the Use-After-Free Error

```
#include <stdlib.h>
```

```
struct node {  
    struct node *next;  
};
```

```
void func(struct node *head) {  
    struct node *p;  
    for (p = head; p != NULL; p = p->next) {  
        free(p);  
    }  
}
```

Prevent Use After Free

- Difficult to detect because these often occur in complex runtime states
 - Allocate in one function
 - Free in another function
 - Use in a third function
- It is not fun to check source code for all possible pointers
 - Are all uses accessing valid (not freed) references?
 - In all possible runtime states

Prevent Use After Free

- What can you do that is not too complex?
 - You can set all freed pointers to NULL
 - Getting a null-pointer dereference if using it
 - Nowadays, OS has built-in defense for null-pointer dereference
 - Then, no one can use them after they are freed
 - Complexity: need to set all aliased pointers to NULL

Related Problem: Double Free

```
main(int argc, char **argv)
{
    ...
    buf1 = (char *) malloc(BUFSIZE1);
    free(buf1);
    buf2 = (char *) malloc(BUFSIZE2);
    strncpy(buf2, argv[1], BUFSIZE2-1);
    free(buf1);
    free(buf2);
}
```

What happens here?

Double Free

- Free buf1, then allocate buf2
 - buf2 may occupy the same memory space of buf1
- buf2 gets user-supplied data

```
strncpy(buf2, argv[1], BUFSIZE2-1);
```
- Free buf1 again
 - Which may use some buf2 data as metadata
 - And may mess up buf2's metadata
- Then free buf2, which uses really messed up metadata

What's Wrong? Fix?

```
#include <stdlib.h>

int f(size_t n) {
    int error_condition = 0;

    int *x = (int *)malloc(n * sizeof(int));
    if (x == NULL)
        return -1;

    /* Use x and set error_condition on error. */
    ...

    if (error_condition == 1) {
        /* Handle error */
        free(x);
    }

    free(x);
    return error_condition;
}
```

What's Wrong? Fix?

```
#include <stdlib.h>
```

```
/* p is a pointer to dynamically allocated memory. */
```

```
void func(void *p, size_t size) {
```

```
    p2 = realloc(p, size);
```

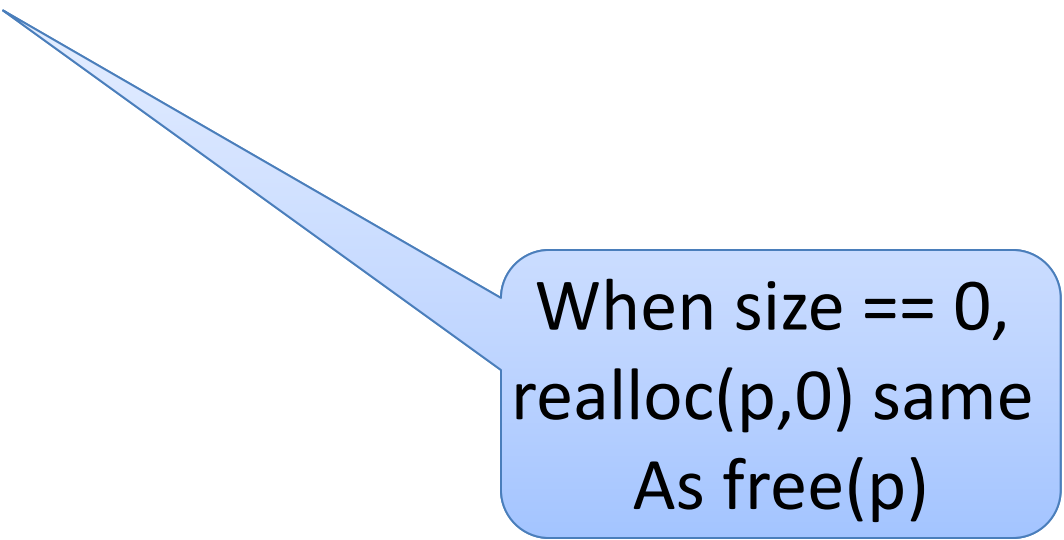
```
    if (p2 == NULL) {
```

```
        free(p);
```

```
        return;
```

```
    }
```

```
}
```



When size == 0,
realloc(p,0) same
As free(p)

Double Free

- So, “double free” can achieve the same effect as some heap overflow vulnerabilities
 - So, can be addressed in the same way
 - But, you can also save yourself some headache by setting freed pointers to NULL
 - Some new heap allocators nowadays have built-in defense against double free

Format String Attacks

Format String Attacks

- Public since 1999
 - first thought as harmless programming errors
- Format string refers to the argument that specifies the format of a string to functions like printf
- Example

```
int i;  
printf ("i = %d with address %08x \n", i, &i);
```



format string

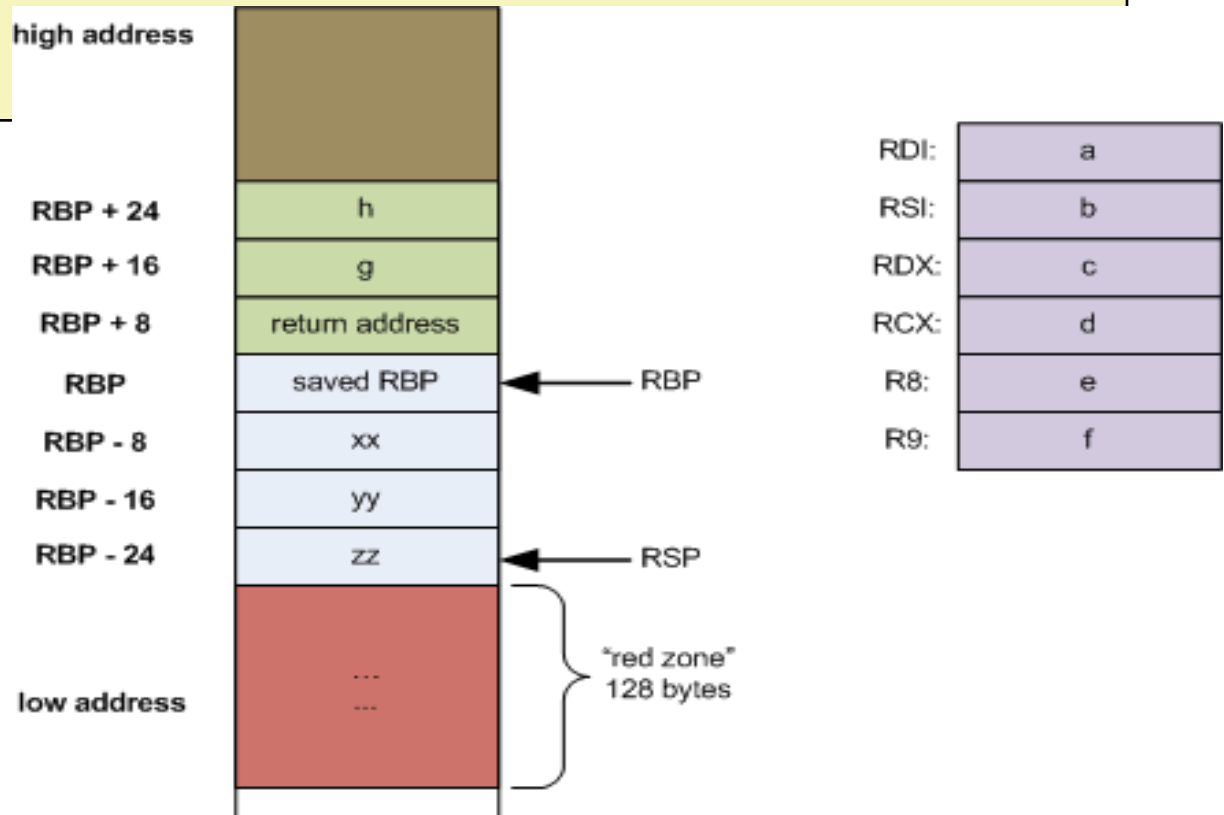
Argument Passing in x86-64

- Arguments (up to the first six) are passed to procedures via registers
 - The rest passed through the stack

Operand size (bits)	Argument Number					
	1	2	3	4	5	6
64	%rdi	%rsi	%rdx	%rcx	%r8	%r9
32	%edi	%esi	%edx	%ecx	%r8d	%r9d
16	%di	%si	%dx	%cx	%r8w	%r9w
8	%dil	%sil	%dl	%cl	%r8b	%r9b

Example of Argument Passing

```
long myfunc(long a, long b, long c, long d,  
            long e, long f, long g, long h) {  
    long xx = a * b * c * d * e * f * g * h;  
    long yy = a + b + c + d + e + f + g + h;  
    long zz = utilfunc(xx, yy, xx % yy);  
    return zz + 20;  
}
```




* Example from <https://eli.thegreenplace.net/2011/09/06/stack-frame-layout-on-x86-64/>

How Does printf Work in C?

```
int i;  
printf ("i = %d with address %08x\n", i, &i);
```

- Pass string addr pointer, i, &i through registers rdi, rsi, rdx, and invokes printf
- When control is inside printf, the function looks for arguments in those registers
 - Or on the stack if more than 6 arguments




How Does printf Work in C?

- What happens for the following program?

```
int i;
```

```
printf ("i = %d with  
address %08x\n");
```

- The C compiler would not complain
 - pretending that the required arguments were at the right place



Format String Attacks

- What about the following simple program for echoing user input?

```
int main(int argc, char *argv[]) {  
    if (argc>1) printf(argv[1]);  
}
```

- Appears to be safe
- However, what would happen if the input is
 - hello%d%d%d%d%d%d
 - Essentially, it runs `printf("hello%d%d%d%d%d%d")`
 - `printf` with only one format-string argument
 - It would print numbers from argument-passing regs and the stack
 - An attacker can view memory this way
- What if the `arg[1]` is "hello%s"?
 - Likely segmentation fault (denial of service)

Format String Attacks

- Getting fancier
 - There is a “%n” specifier, which writes the number of bytes already printed, into a variable of our choice.

```
int i;
```

```
printf ("foobar%n\n", (int *) &i);
```

```
printf ("i = %d\n", i);
```

- i gets 6 at the end

Format String Attacks

- For the echo program, what if the user input is “foobar%n”?
 - It will take an address from rsi, and write 6 to the memory slot with that address
 - What about “foobar%10u%n”?
 - Write possibly 16 to a memory location
 - How to write to an arbitrary address?
 - Put that address at the right place (in the right register or stack slot)
- Therefore, an attacker can possibly update any memory with arbitrary contents
 - How about overwriting a function pointer and hijacking the control flow (and installing some worm code)?

Format String Attacks

```
int main(int argc, char *argv[]) {  
    char buf[512];  
    fgets(buf, sizeof(buf), stdin);
```



No buffer
overflow
here

```
    printf("The input is:");  
    printf(buf);  
    return 0;
```



But, format
string attacks

- ```
}
```
- An attacker can
    - view/change any part of the memory
    - execute arbitrary code
      - Just put the code into buf
  - More details see paper
    - “Exploiting Format String Vulnerabilities”

# Format String Attacks: Fixes

---

- Most of time: quite easy to fix

```
int main(int argc, char *argv[]) {
 printf(argv[1]); printf("%s", argv[1])
 printf("\n");
}
```

- But not always so obvious
  - sometimes not easy to find

```
void foo(char *user) {
 char outbuf[512];
 char buffer[512];
 sprintf (buffer, "ERR Wrong command: %400s", user);
 sprintf (outbuf, buffer);
}
```

Is there a buffer overflow? Is there a format-string vulnerability?  
How to fix it?

# Prevent Format String Vulnerabilities

---

- Preventing format string vulnerabilities means limiting the ability of adversaries to control the format string
  - Hard-coded strings w/ no arguments – when you can
  - Hard-coded format strings at least – no `printf(arg)`
  - Do not use `%n`
    - Be careful with other references - `%s` and `sprintf` can be used to create disclosure attacks
  - Compiler support to match `printf` arguments with format string



# Summary: Memory-Corruption Vulnerabilities

---

- Buffer overflow
  - stack smashing: overwrite data on the stack
  - can also overwrite data on the heap
  - Integer overflow makes it easier
- Use after free; double free
- Format string attacks
- And there are many more ...
  - E.g., Mark Dowd showed it was possible to hijack Adobe Flash player based on a [null pointer dereference](#)
- Type-safe languages have other kinds of vulnerabilities
  - E.g., SQL injection attacks, ...